

Designing Efficient AI Model

Chia-Chi Tsai (蔡家齊)

cctsai@gs.ncku.edu.tw

AI System Lab

Department of Electrical Engineering

National Cheng Kung University

Outline



- Basic Concepts
- Manual Designed Neural Networks


Outline



- Basic Concepts
- Manual Designed Neural Networks

Number of MAC Operations

- Linear Layer



Layer Type	MACs (batch size n=1)
Linear Layer	$c_o \times c_i$
Convolution	
Grouped Convolution	
Depthwise Convolution	
1x1 Convolution	



Notation	
n	Batch size
c_i, c_o	Input/output channel
w_i, w_o	Input/output width
h_i, h_o	Input/output height
k_w, k_h	Kernel width/height
g	Groups

Number of MAC Operations



- Convolution

Layer Type	MACs (batch size n=1)
Linear Layer	$c_o \times c_i$
Convolution	$c_i \times k_h \times k_w \times h_o \times w_o \times c_o$
Grouped Convolution	
Depthwise Convolution	
1x1 Convolution	




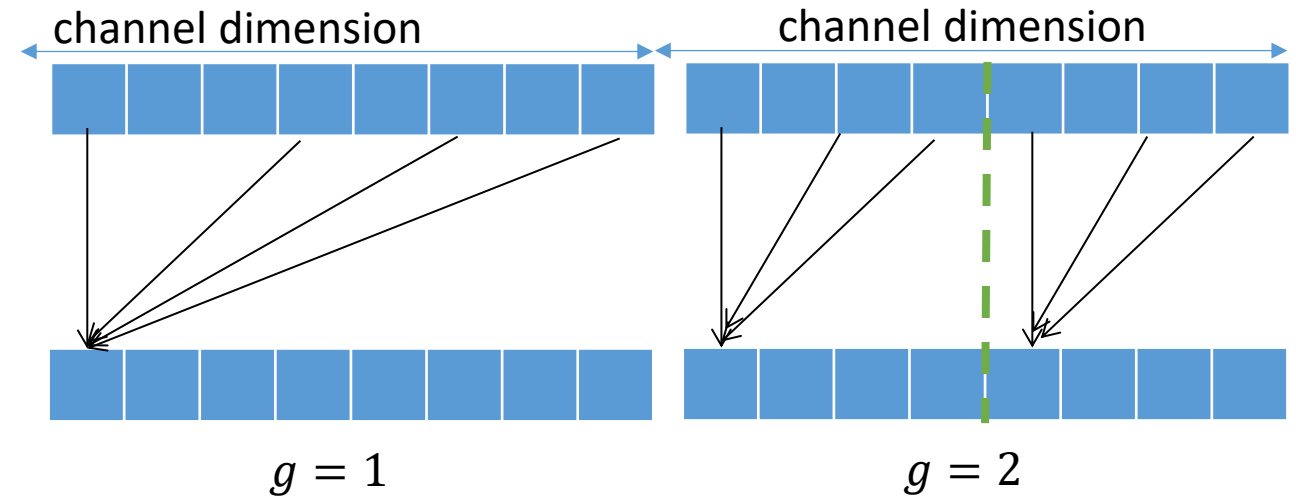
Notation	
n	Batch size
c_i, c_o	Input/output channel
w_i, w_o	Input/output width
h_i, h_o	Input/output height
k_w, k_h	Kernel width/height
g	Groups

Number of MAC Operations



- Grouped Convolution

Layer Type	MACs (batch size n=1)
Linear Layer	$c_o \times c_i$
Convolution	$c_i \times k_h \times k_w \times h_o \times w_o \times c_o$
Grouped Convolution	$\frac{c_i}{g} \times k_h \times k_w \times h_o \times w_o \times c_o$
Depthwise Convolution	
1x1 Convolution	





Notation	
n	Batch size
c_i, c_o	Input/output channel
w_i, w_o	Input/output width
h_i, h_o	Input/output height
k_w, k_h	Kernel width/height
g	Groups

Number of MAC Operations



- Depthwise Convolution

Layer Type	MACs (batch size n=1)
Linear Layer	$c_o \times c_i$
Convolution	$c_i \times k_h \times k_w \times h_o \times w_o \times c_o$
Grouped Convolution	$\frac{c_i}{g} \times k_h \times k_w \times h_o \times w_o \times c_o$
Depthwise Convolution	$k_h \times k_w \times h_o \times w_o \times c_o$
1x1 Convolution	





Notation	
n	Batch size
c_i, c_o	Input/output channel
w_i, w_o	Input/output width
h_i, h_o	Input/output height
k_w, k_h	Kernel width/height
g	Groups

Number of MAC Operations



- Convolution

Layer Type	MACs (batch size n=1)
Linear Layer	$c_o \times c_i$
Convolution	$c_i \times k_h \times k_w \times h_o \times w_o \times c_o$
Grouped Convolution	$\frac{c_i}{g} \times k_h \times k_w \times h_o \times w_o \times c_o$
Depthwise Convolution	$k_h \times k_w \times h_o \times w_o \times c_o$
1x1 Convolution	$c_o \times c_i \times h_o \times w_o$



Notation	
n	Batch size
c_i, c_o	Input/output channel
w_i, w_o	Input/output width
h_i, h_o	Input/output height
k_w, k_h	Kernel width/height
g	Groups

Stages of Model

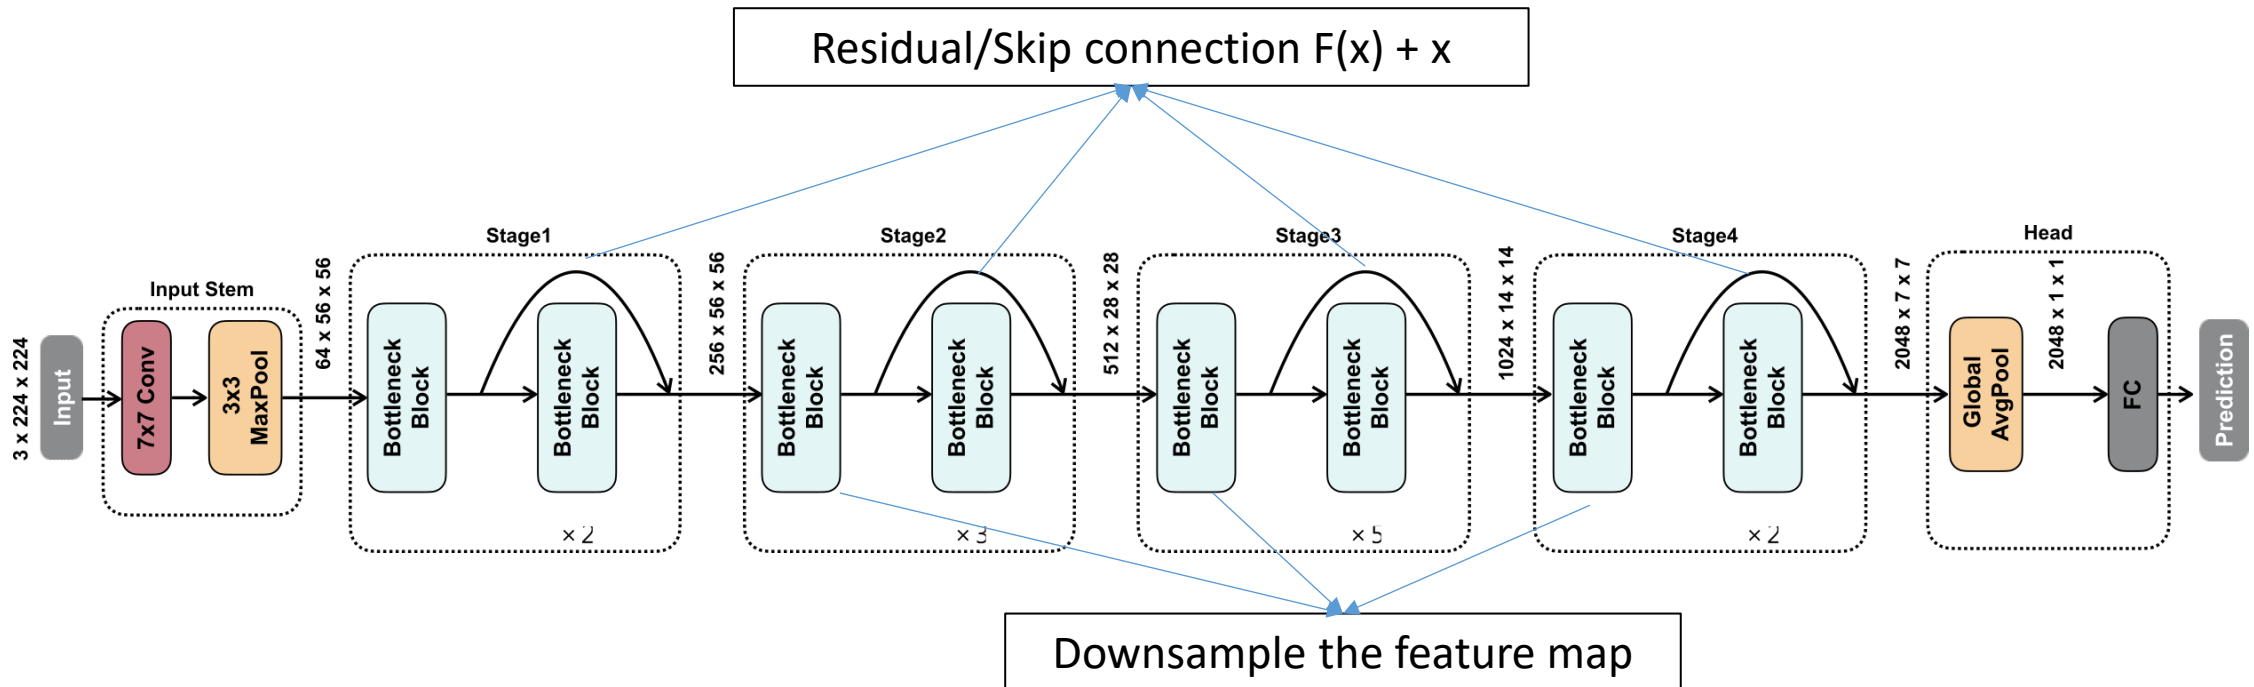


- A neural network architecture typically consists
 - Input stem,
 - Head
 - Several stages
- Early stages have larger feature map sizes, so we need to keep the width small to reduce the cost.
- In contrast, late stages have smaller feature map sizes so that we can increase the width.

Stages of Model





- Downsample
 - Feature map downsampling is usually done at the first block in each stage via stride convolution or pooling
- Residual/Skip connection
 - We can add residual/skip connections for the remaining blocks as their input and output
 - dimensions are the same



Basic Concepts



- Residual/Skip connection



Outline



- Basic Concepts
- Manual Designed Neural Networks


AlexNet



- Remarkable improvements over previous non-DL methods

Image (3×224×224)
11×11 Conv, channel 96, stride 4, pad 2
3×3 MaxPool, stride 2
5×5 Conv, channel 256, pad 2, groups 2
3×3 MaxPool, stride 2
3×3 Conv, channel 384, pad 1
3×3 Conv, channel 384, pad 1, groups 2
3×3 Conv, channel 256, pad 1, groups 2
3×3 MaxPool, stride 2
Linear, channel 4096
Linear, channel 4096
Linear, channel 1000

Large kernel convolution in early stages




Non-DL method

DL method


Image (3×224×224)
3×3 Conv, channel 64, pad 1
3×3 Conv, channel 64, pad 1
2×2 MaxPool, stride 2
3×3 Conv, channel 128, pad 1
3×3 Conv, channel 128, pad 1
2×2 MaxPool, stride 2
3×3 Conv, channel 256, pad 1
3×3 Conv, channel 256, pad 1
3×3 Conv, channel 256, pad 1
2×2 MaxPool, stride 2
3×3 Conv, channel 512, pad 1
3×3 Conv, channel 512, pad 1
3×3 Conv, channel 512, pad 1
2×2 MaxPool, stride 2
3×3 Conv, channel 512, pad 1
3×3 Conv, channel 512, pad 1
3×3 Conv, channel 512, pad 1
2×2 MaxPool, stride 2
Linear, channel 4096
Linear, channel 4096
Linear, channel 1000

- Stacking multiple 3x3 convolution layers
- Different from AlexNet, VGG only uses 3x3 convolution. Meanwhile, VGG
- Stacks more layers to maintain a large receptive field.
 - The computational cost of two 3x3 convolutions is smaller than one 5x5 convolution: $3 \times 3 + 3 \times 3 = 18 < 5 \times 5$



- VGG: performance and cost breakdown

ImageNet Classification, top-5 error (%)




Layer	Weights	FLOP
Conv1_1	2K	0.2B
Conv1_2	37K	3.7B
Conv2_1	74K	1.8B
Conv2_2	148K	3.7B
Conv3_1	295K	1.8B
Conv3_2	590K	3.7B
Conv3_3	590K	3.7B
Conv4_1	1M	1.8B
Conv4_2	2M	3.7B
Conv4_3	2M	3.7B
Conv5_1	2M	925M
Conv5_2	2M	925M
Conv5_3	2M	925M
Fc6	103M	206M
Fc7	17M	34M
Fc8	4M	8M
Total	138M	30.9B

3x3 convolution layers are the key computational bottleneck

SqueezeNet



- Replace 3x3 convolution with fire modules
- Use global average pooling in the head to reduce the cost of the head

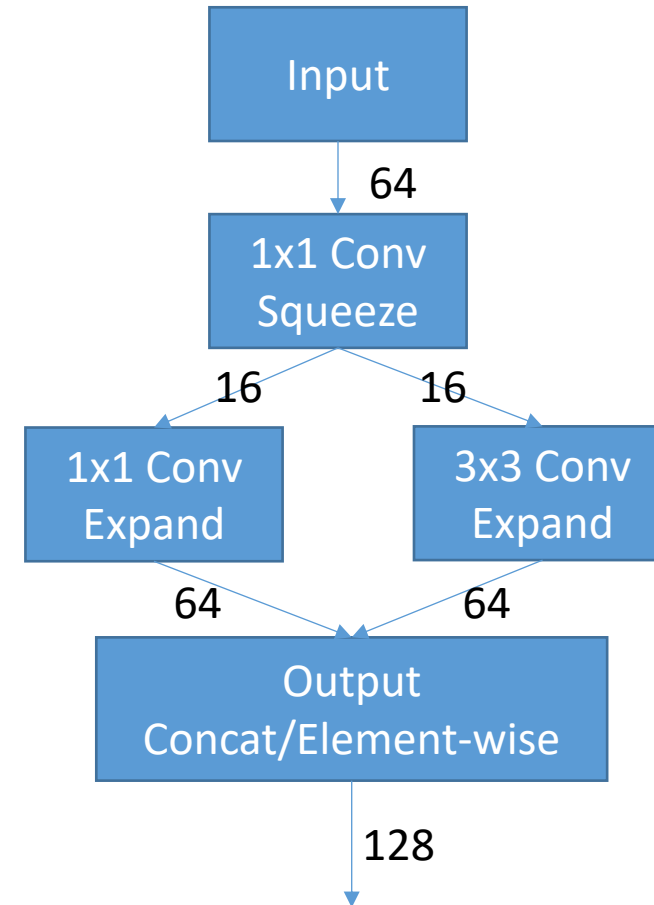


landola, F. N., Han, S., Moskewicz, M. W., Ashraf, K., Dally, W. J., & Keutzer, K. (2016). SqueezeNet: AlexNet-level accuracy with 50x fewer parameters and < 0.5 MB model size. *arXiv preprint arXiv:1602.07360*.

Fire Module



- Reduce the cost by using 1x1 convolution
 - Replace some 3x3 filters with 1x1 filters
 - Decrease the number of input channels to 3x3 filters.



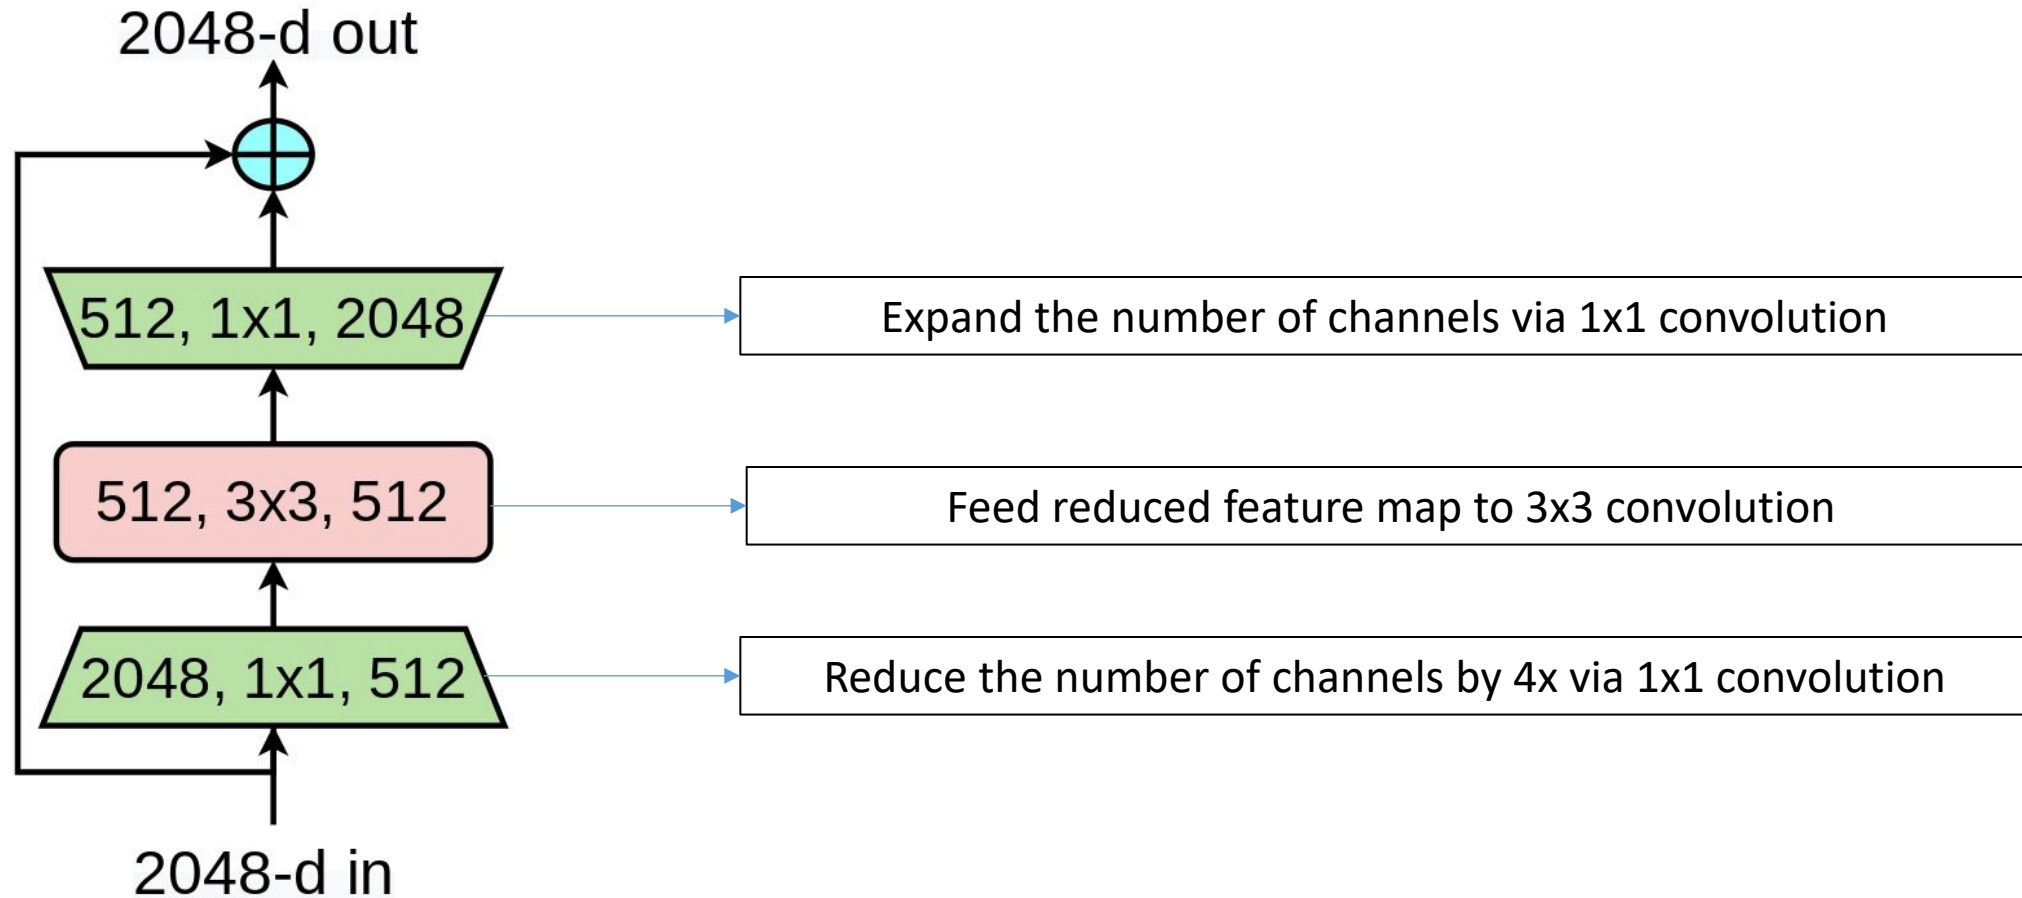
SqueezeNet



- AlexNet-level accuracy with 50x-510x smaller model size

CNN architecture	Compression Approach	Data Type	Original → Compressed Model Size	Reduction in Model Size vs. AlexNet	Top-1 ImageNet Accuracy	Top-5 ImageNet Accuracy
AlexNet	None (baseline)	32 bit	240MB	1x	57.2%	80.3%
AlexNet	SVD (Denton et al., 2014)	32 bit	240MB → 48MB	5x	56.0%	79.4%
AlexNet	Network Pruning (Han et al., 2015b)	32 bit	240MB → 27MB	9x	57.2%	80.3%
AlexNet	Deep Compression (Han et al., 2015a)	5-8 bit	240MB → 6.9MB	35x	57.2%	80.3%
SqueezeNet (ours)	None	32 bit	4.8MB	50x	57.5%	80.3%
SqueezeNet (ours)	Deep Compression	8 bit	4.8MB → 0.66MB	363x	57.5%	80.3%
SqueezeNet (ours)	Deep Compression	6 bit	4.8MB → 0.47MB	510x	57.5%	80.3%

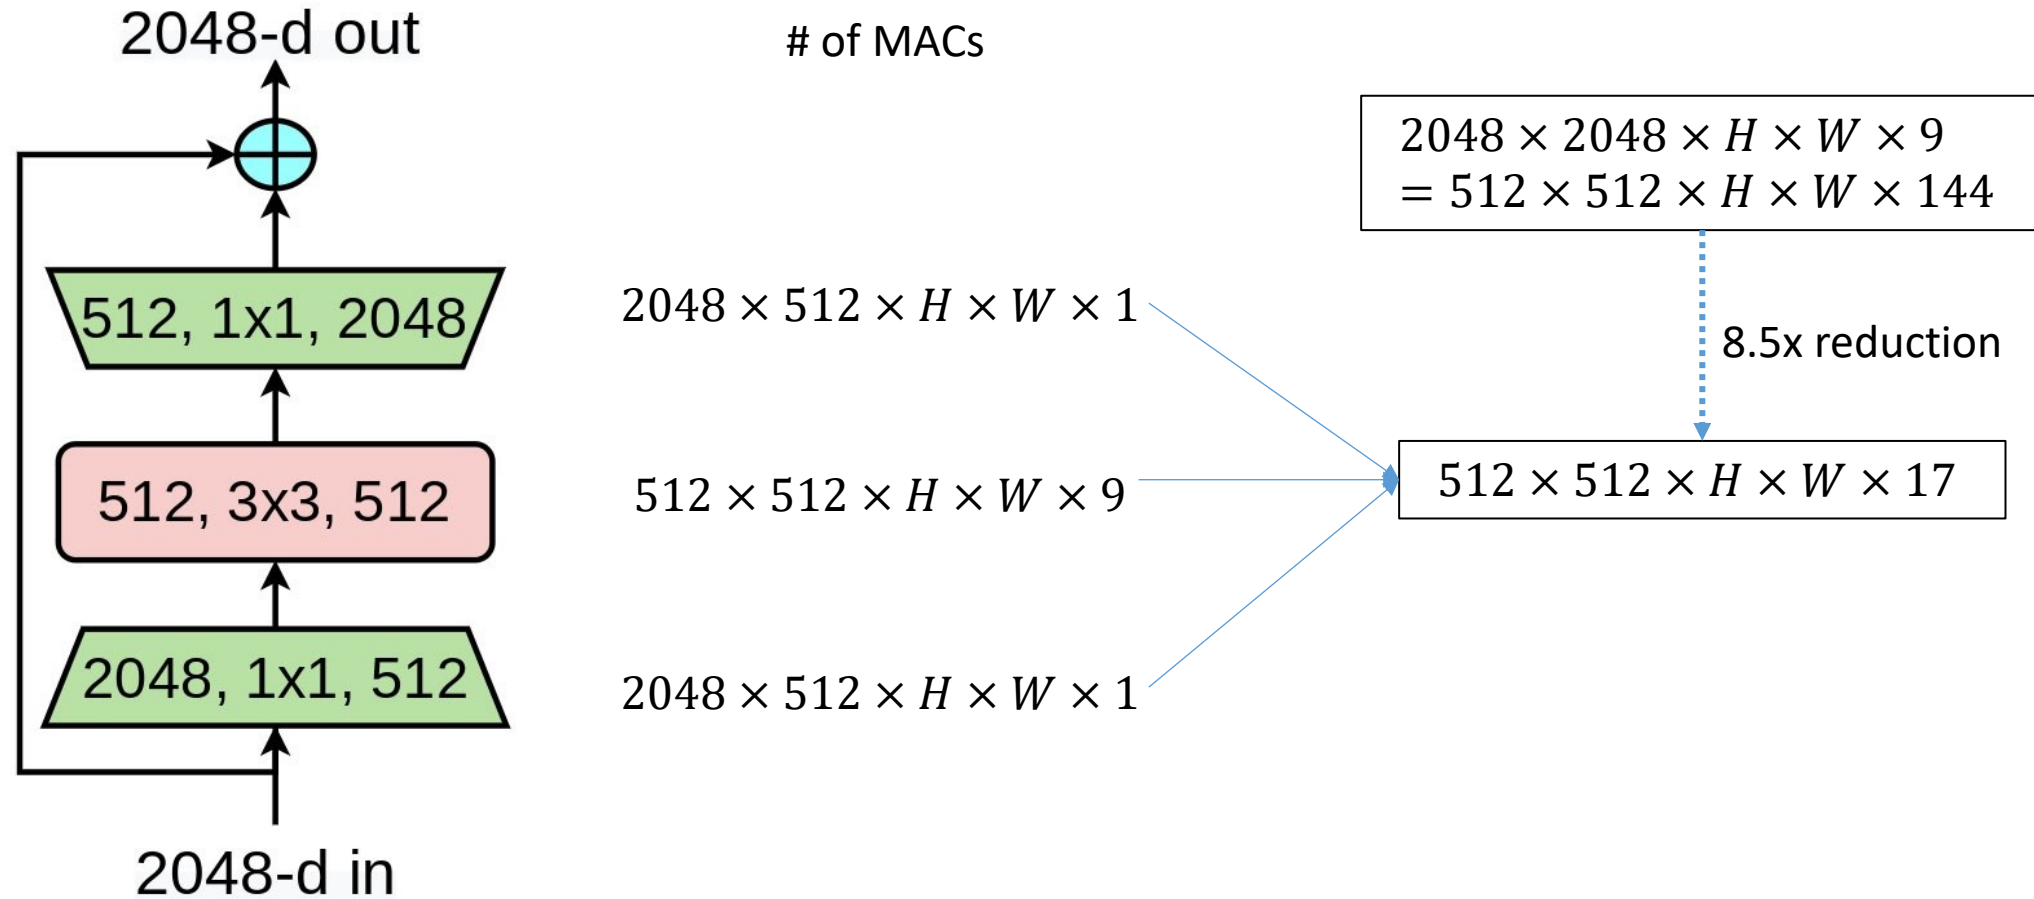
ResNet50: Bottleneck Block



ResNet Bottleneck

He, K., Zhang, X., Ren, S., & Sun, J. (2016). Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition* (pp. 770-778).

ResNet50: Bottleneck Block

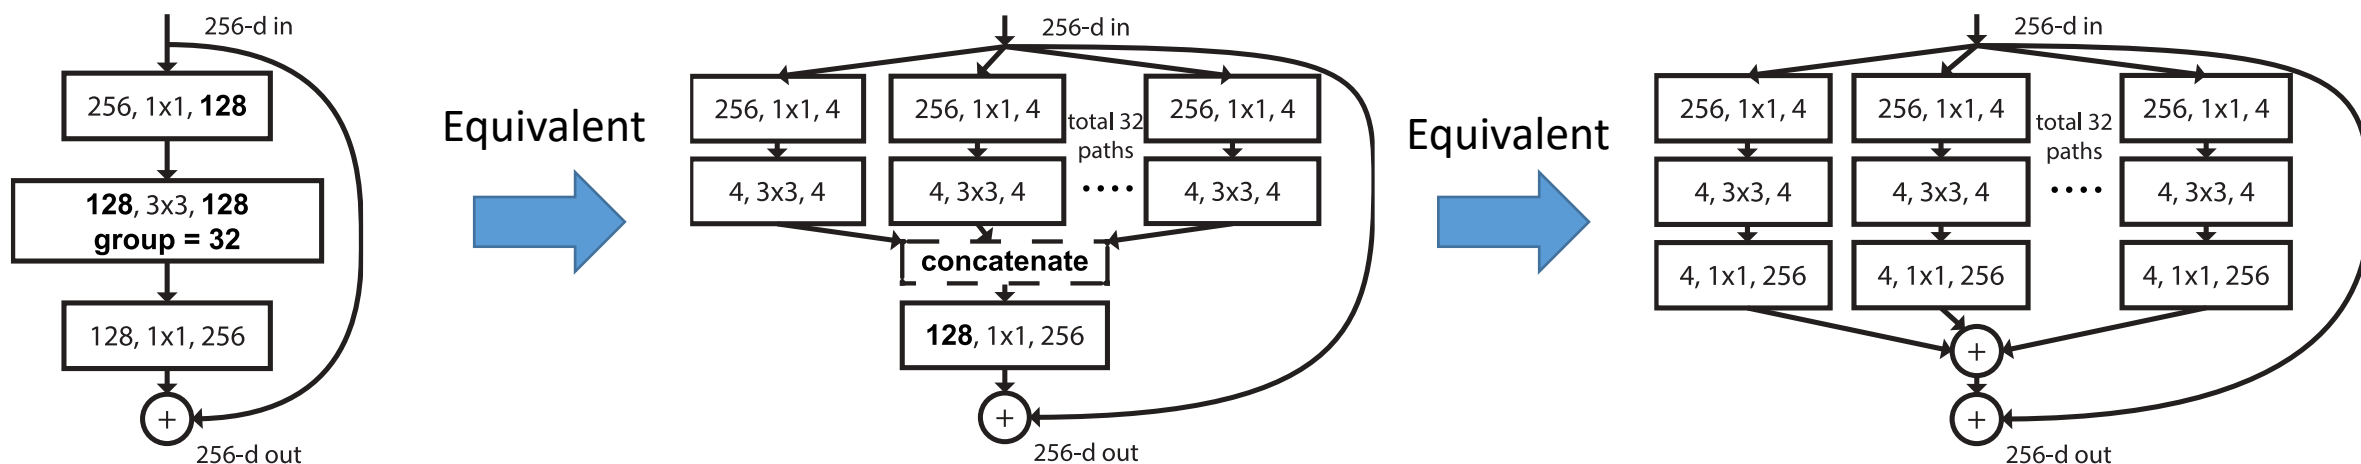


ResNet Bottleneck

ResNeXt: Grouped Convolution



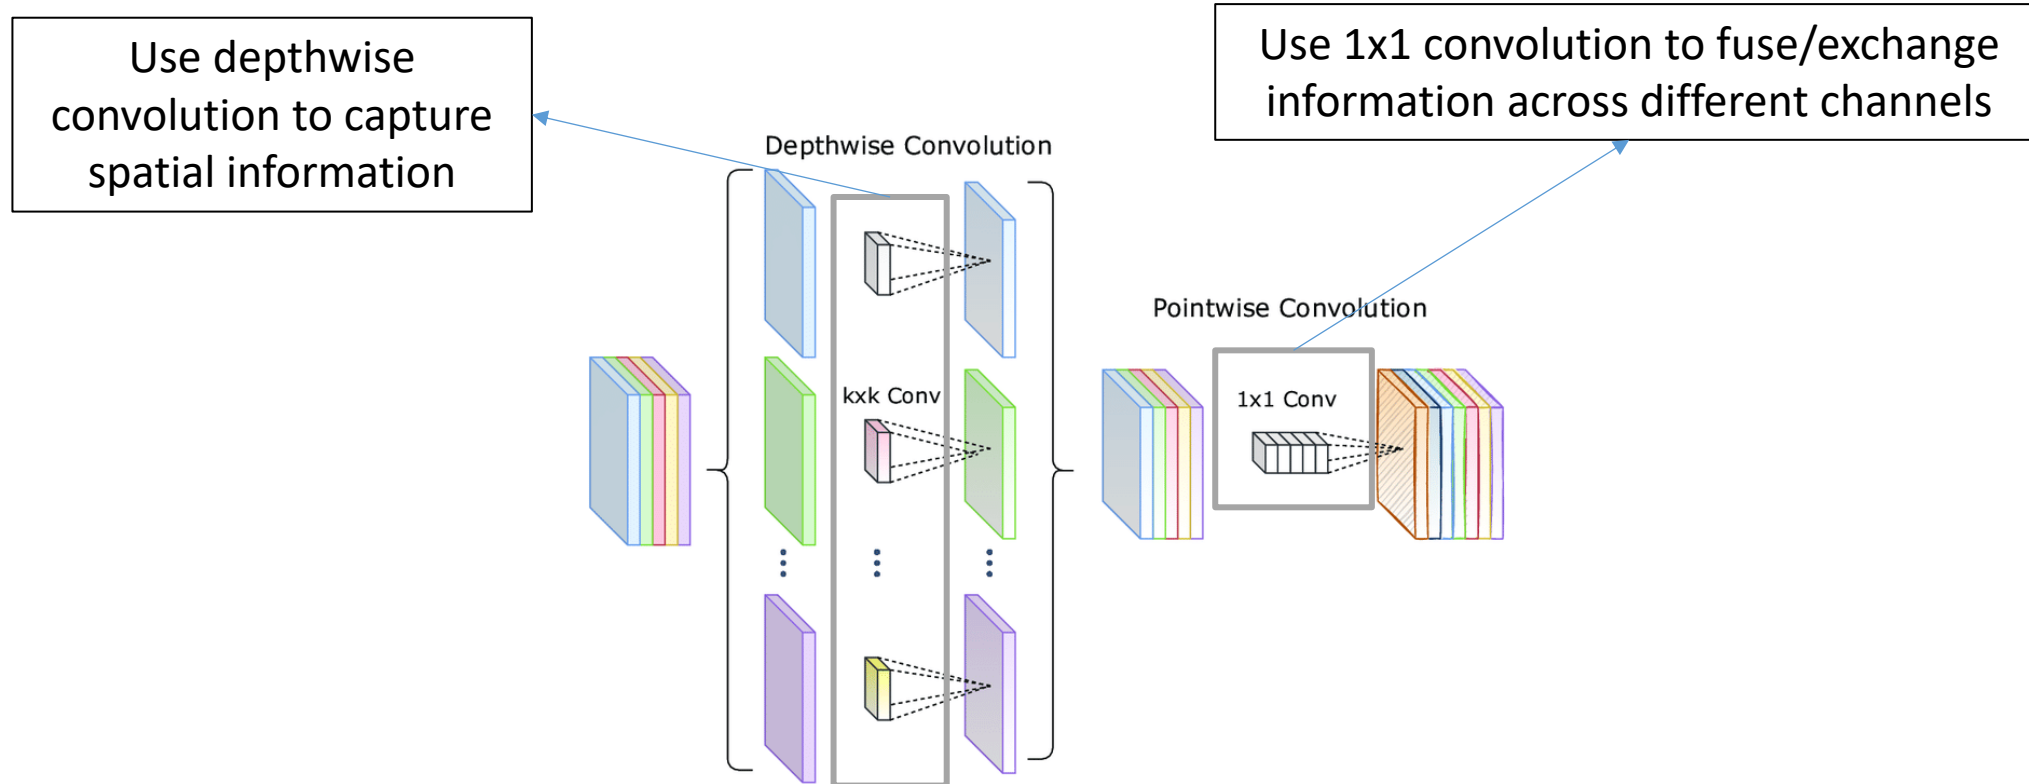
- Replace 3x3 convolution with 3x3 grouped convolution
- Equivalent to a multi-path block



MobileNet: Depthwise-Separable Block



- Depthwise convolution - An extreme case of group convolution
 - Where the group number equals the number of input channels




MobileNetV2: Inverted Bottleneck Block




- Depthwise convolution has a much lower capacity compared to normal convolution
 - Increase the depthwise convolution's input and output channels to improve its capacity
 - Depthwise convolution's cost only grows linearly. Therefore, the cost is still affordable.

MobileNetV1



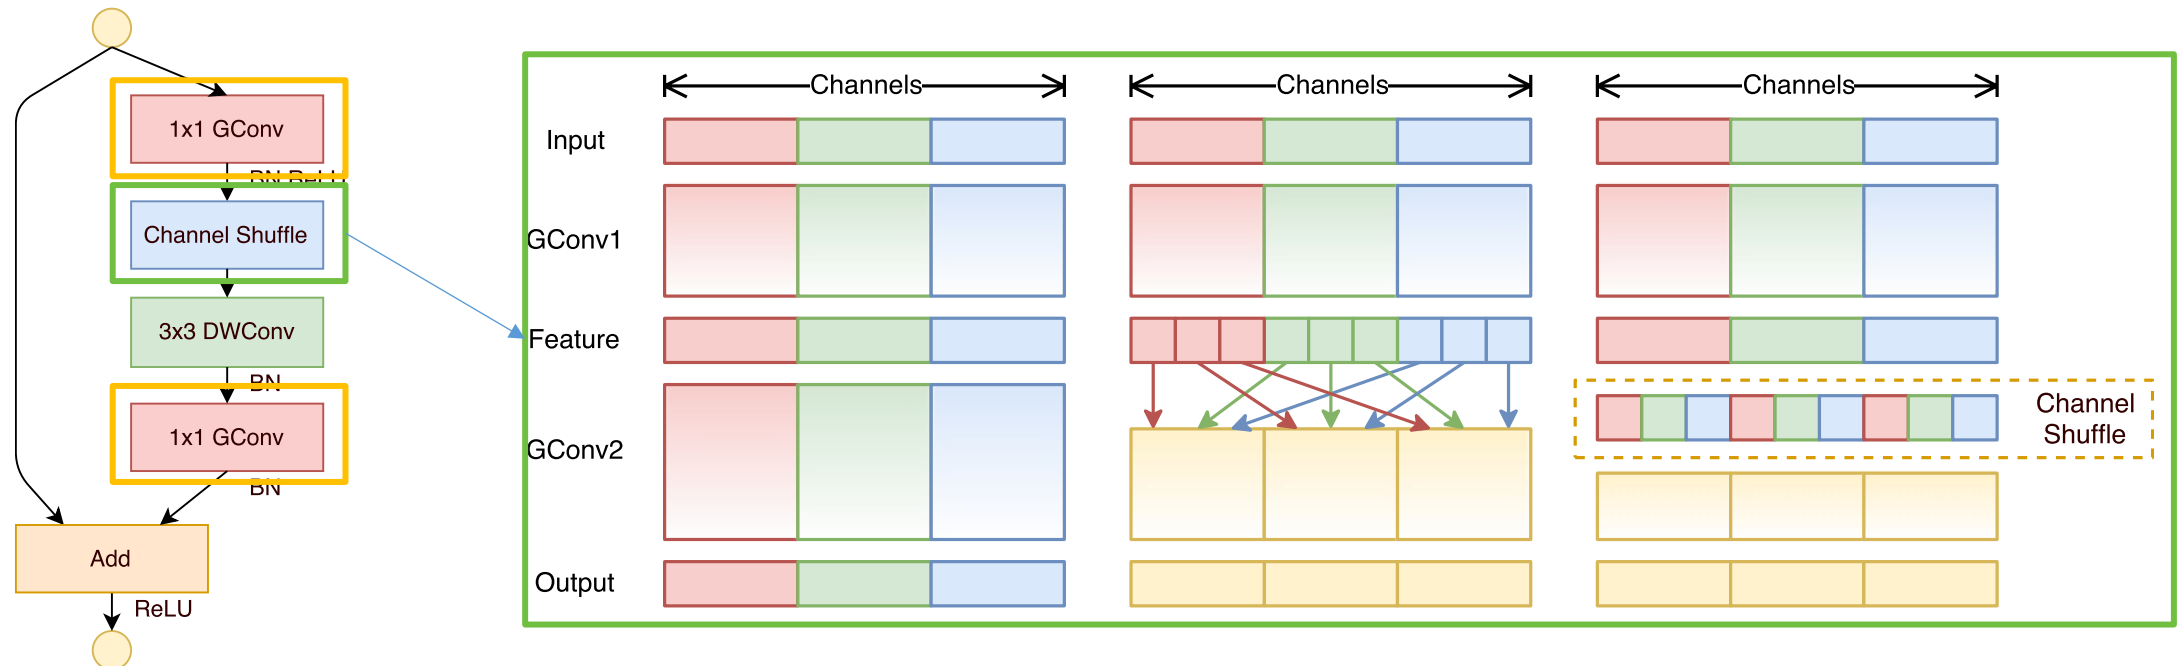
MobileNetV2



ShuffleNet



- 1x1 Group Convolution
 - Further reduce the cost by replacing 1x1 convolution with 1x1 group convolution
- Channel Shuffle
 - Exchange information across different groups via channel shuffle




Zhang, X., Zhou, X., Lin, M., & Sun, J. (2018). Shufflenet: An extremely efficient convolutional neural network for mobile devices. In *Proceedings of the IEEE conference on computer vision and pattern recognition* (pp. 6848-6856).

Accuracy-Efficiency Trade-off



- Huge design space, manual design is unscalable



From Manual Design to Automatic Design



- Automatic Architecture Search

